



B.S. Abdur Rahman
Crescent
Institute of Science & Technology
Deemed to be University u/s 3 of the UGC Act, 1956

DEPARTMENT OF COMPUTER APPLICATIONS

CADX116 - MACHINE LEARNING ALGORITHMS

LABORATORY



B.S. Abdur Rahman

Crescent

Institute of Science & Technology

Deemed to be University u/s 3 of the UGC Act, 1956

LAB RECORD

NAME : _____

RRN : _____

LAB : CADX116 - MACHINE LEARNING ALGORITHMS LABORATORY



B.S. Abdur Rahman
Crescent
Institute of Science & Technology
Deemed to be University u/s 3 of the UGC Act, 1956

DEPARTMENT OF COMPUTER APPLICATIONS

ACADEMIC YEAR (JULY 2024 - DECEMBER 2024)

COURSE CODE : CADX116

***COURSE NAME : MACHINE LEARNING ALGORITHMS
LABORATORY***

PROGRAMME : BCA (DATA SCIENCE)

SEMESTER : V



B.S. Abdur Rahman
Crescent
Institute of Science & Technology
Deemed to be University u/s 3 of the UGC Act, 1956

BONAFIDE CERTIFICATE

This is a Certified Record Book of _____

RRN: _____ submitted for the Semester End

Practical Examination held on _____, for the CADX116 -

MACHINE LEARNING ALGORITHMS LABORATORY during 2024 - 2025.

.....
Signature of Faculty

INDEX

<i>Ex.No.</i>	<i>Date</i>	<i>List of Experiments</i>	<i>Page No</i>	<i>Signature</i>
<i>1.</i>		<i>Implement Naive Bayes Classifier</i>		
<i>2.</i>		<i>Implement Linear Regression</i>		
<i>3.</i>		<i>Implement Support Vector Machine (SVM) Algorithm</i>		
<i>4.</i>		<i>Implement Decision Tree Classifier</i>		
<i>5.</i>		<i>Implement Random Forest Algorithm</i>		
<i>6.</i>		<i>Implement K-Means Clustering Algorithm</i>		
<i>7.</i>		<i>Implement Principal Component Analysis (PCA)</i>		
<i>8.</i>		<i>Implement K-Nearest Neighbors (KNN) Algorithm</i>		
<i>9.</i>		<i>Demonstrate Weka Tool</i>		
<i>10.</i>		<i>Build An Unsupervised Models Using [Scikit-Learn And Pytorch]</i>		
<i>11.</i>		<i>Implement Logistic Regression</i>		
<i>12.</i>		<i>Implement Neural Network Model</i>		

EX.NO :

DATE :

IMPLEMENT NAÏVE BAYES CLASSIFIER

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
    from sklearn import tree
    from sklearn.preprocessing import LabelEncoder
    from sklearn.naive_bayes import GaussianNB
```

Step 2:

```
[ ] data=pd.read_csv("Naive_Bayes_PlayTennis.csv")
    print("the First 5 values of data is :\n")
    data.head()
```

Output:

the First 5 values of data is :

	Outlook	Temperature	Humidity	Windy	Play Tennis
0	<i>Sunny</i>	<i>Hot</i>	<i>High</i>	<i>False</i>	<i>No</i>
1	<i>Sunny</i>	<i>Hot</i>	<i>High</i>	<i>True</i>	<i>No</i>
2	<i>Overcast</i>	<i>Hot</i>	<i>High</i>	<i>False</i>	<i>Yes</i>
3	<i>Rainy</i>	<i>Mild</i>	<i>High</i>	<i>False</i>	<i>Yes</i>
4	<i>Rainy</i>	<i>Cool</i>	<i>Normal</i>	<i>False</i>	<i>Yes</i>

Step 3:

```
[ ] X=data.iloc[:, :-1]
    print("the First 5 values of data is :\n")
    X.head()
```

Output:

the First 5 values of data is :

	<i>Outlook</i>	<i>Temperature</i>	<i>Humidity</i>	<i>Windy</i>
<i>0</i>	<i>Sunny</i>	<i>Hot</i>	<i>High</i>	<i>False</i>
<i>1</i>	<i>Sunny</i>	<i>Hot</i>	<i>High</i>	<i>True</i>
<i>2</i>	<i>Overcast</i>	<i>Hot</i>	<i>High</i>	<i>False</i>
<i>3</i>	<i>Rainy</i>	<i>Mild</i>	<i>High</i>	<i>False</i>
<i>4</i>	<i>Rainy</i>	<i>Cool</i>	<i>Normal</i>	<i>False</i>

Step 4:

```
[ ] y=data.iloc[:,-1]
    print("the First 5 values of data is :\n")
    y.head()
```

Output:

the First 5 values of data is :

```
0    No
1    No
2    Yes
3    Yes
4    Yes
Name: PlayTennis, dtype: object
```

Step 5:

```
[ ] for column in X.columns:
    X[column]=LabelEncoder().fit_transform(X[column])
    X.head()
```

Output:

	<i>Outlook</i>	<i>Temperature</i>	<i>Humidity</i>	<i>Windy</i>
<i>0</i>	<i>2</i>	<i>1</i>	<i>0</i>	<i>0</i>
<i>1</i>	<i>2</i>	<i>1</i>	<i>0</i>	<i>1</i>
<i>2</i>	<i>0</i>	<i>1</i>	<i>0</i>	<i>0</i>
<i>3</i>	<i>1</i>	<i>2</i>	<i>0</i>	<i>0</i>
<i>4</i>	<i>1</i>	<i>0</i>	<i>1</i>	<i>0</i>

Step 6:

```
[ ] y=LabelEncoder().fit_transform(y)
    print(y)
```

Output:

[0 0 1 1 1 0 1 0 1 1 1 1 1 0]

Step 7:

```
[ ] from sklearn.model_selection import train_test_split
```

Step 8:

```
[ ] X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.20)
```

Step 9:

```
[ ] from sklearn.naive_bayes import GaussianNB
    classifier = GaussianNB()
    classifier.fit(X_train,y_train)
```

Output:

```
▼ GaussianNB
GaussianNB()
```

Step 10:

```
[ ] from sklearn.metrics import accuracy_score
```

Step 11:

```
[ ] print("Accuracy is:",accuracy_score(classifier.predict(X_test),(y_test)))
```

Output:

Accuracy is: 0.6666666666666666

EX.NO :

DATE :

IMPLEMENT LINEAR REGRESSION

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

Step 2:

```
[ ] data=pd.read_csv("Linear_Regression_Advertising.csv")
print("the First 5 values of data is :\n")
data.head()
```

Output:

the First 5 values of data is :

	<i>TV</i>	<i>Radio</i>	<i>Newspaper</i>	<i>Sales</i>
<i>0</i>	230.1	37.8	69.2	22.1
<i>1</i>	44.5	39.3	45.1	10.4
<i>2</i>	17.2	45.9	69.3	9.3
<i>3</i>	151.5	41.3	58.5	18.5
<i>4</i>	180.8	10.8	58.4	12.9

Step 3:

```
[ ] X = data[['TV', 'Radio', 'Newspaper']]
    y = data['Sales']
    X.head()
```

Output:

	<i>TV</i>	<i>Radio</i>	<i>Newspaper</i>
<i>0</i>	230.1	37.8	69.2
<i>1</i>	44.5	39.3	45.1
<i>2</i>	17.2	45.9	69.3
<i>3</i>	151.5	41.3	58.5
<i>4</i>	180.8	10.8	58.4

Step 4:

```
[ ] y.head()
```

Output:

```
0    22.1
1    10.4
2     9.3
3    18.5
4    12.9
Name: Sales, dtype: float64
```

Step 5:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
    model = LinearRegression()
    model.fit(X_train, y_train)
```

Output:

```
▼ LinearRegression
LinearRegression()
```

Step 6:

```
[ ] y_pred = model.predict(X_test)
```

Step 7:

```
[ ] mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    print(f'Mean Squared Error: {mse}')
    print(f'R^2 Score: {r2}')
```

Output :

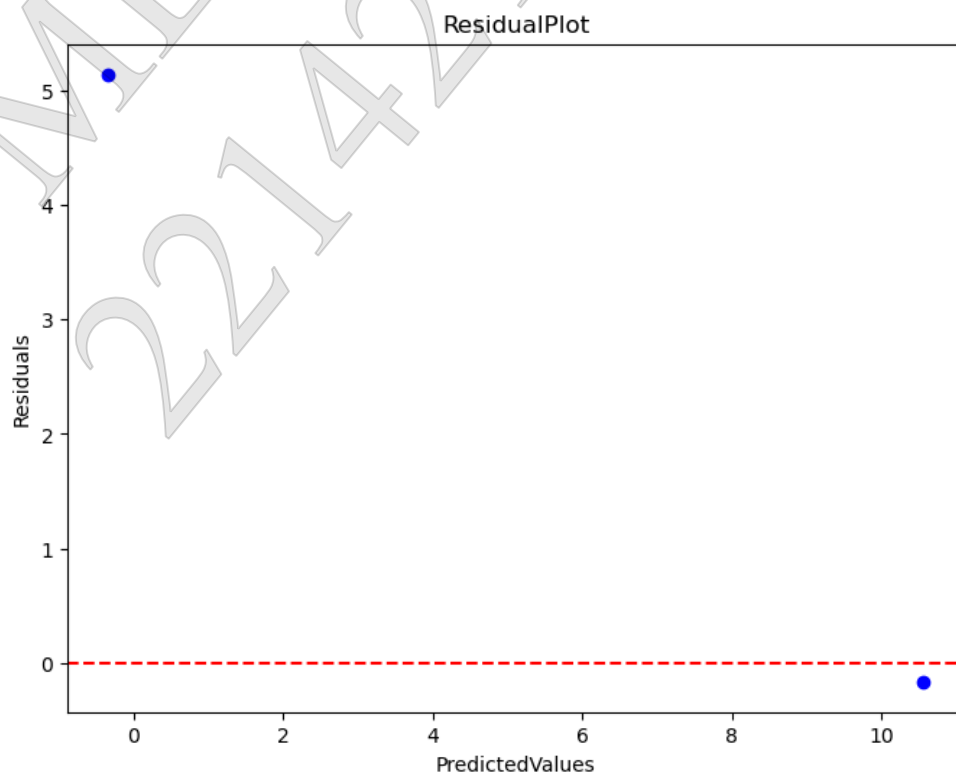
Mean Squared Error: 13.201451790963432

R^2 Score: -0.6838586468065599

Step 8:

```
[ ] residuals=y_test-y_pred
    plt.figure(figsize=(8,6))
    plt.scatter(y_pred,residuals,color='blue')
    plt.title('ResidualPlot')
    plt.xlabel('PredictedValues')
    plt.ylabel('Residuals')
    plt.axhline(y=0,color='r', linestyle='--')
    plt.show()
```


Output:



Step 9:

```
[ ] plt.figure(figsize=(8,6))
    plt.scatter(y_test,y_pred,color='green')
    plt.title('PredictedvsActual')
    plt.xlabel('ActualValues')
    plt.ylabel('PredictedValues')
    plt.plot([min(y_test),max(y_test)], [min(y_test),max(y_test)],color='red',
    linestyle='--')
    plt.show()
```

Output:



EX.NO :

DATE :

IMPLEMENT SUPPORT VECTOR MACHINE (SVM) ALGORITHM

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,
recall_score, f1_score
from sklearn.svm import SVC
```

Step 2:

```
[ ] data=pd.read_csv("SVM_PurchasePrediction.csv")
print("the First 5 values of data is :\n")
data.head()
```

Output:

the First 5 values of data is :

	<i>Age</i>	<i>Income</i>	<i>Gender</i>	<i>Purchase</i>
<i>0</i>	25	50000	Male	0
<i>1</i>	45	64000	Female	1
<i>2</i>	35	58000	Female	0
<i>3</i>	50	72000	Male	1
<i>4</i>	23	48000	Male	0

Step 3:

```
[ ] le=LabelEncoder()  
    data['Gender']=le.fit_transform(data['Gender'])
```

Step 4:

```
[ ] data.fillna(method='ffill',inplace=True)
```

Step 5:

```
[ ] X =data.drop('Purchase',axis=1)  
    y = data['Purchase']
```

Step 6:

```
[ ] sc=StandardScaler()  
    X_scaled=sc.fit_transform(X)
```

Step 7:

```
[ ] classifiers={  
    'Linear SVM': SVC(kernel='linear',random_state=0),  
    'Polynomial SVM': SVC(kernel='poly',degree=3,random_state=0),  
    'RBF SVM': SVC(kernel='rbf',random_state=0)  
}
```

Step 8:

```
[ ] for clf_name,clf in classifiers.items():  
    scores=cross_val_score(clf,X_scaled,y,cv=5,scoring='accuracy')  
    print(f"{clf_name} Cross-validation Accuracy : {scores.mean():.2f}(+/-  
    {scores.std()*2:.2f})")
```

Output:

Linear SVM Cross-validation Accuracy : 0.80(+/-0.33)
Polynomial SVM Cross-validation Accuracy : 0.80(+/-0.53)
RBF SVM Cross-validation Accuracy : 0.80(+/-0.53)

Step 9:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)  
    classifier=SVC(kernel='linear',random_state=0)  
    classifier.fit(X_train, y_train)
```

Output:

```
SVC
SVC(kernel='linear', random_state=0)
```

Step 10:

```
[] y_pred=classifier.predict(X_test)
```

Step 11:

```
[] cm=confusion_matrix(y_test,y_pred)
accuracy=accuracy_score(y_test,y_pred)
precision=precision_score(y_test,y_pred)
recall=recall_score(y_test,y_pred)
f1=f1_score(y_test,y_pred)
```

Step 12:

```
[] print(f"Confusion Matrix :\n",cm)
print(f"Accuracy :{accuracy:.2f}")
print(f"Precision :{precision:.2f}")
print(f"Recall :{recall:.2f}")
print(f"F1 Performance score :{f1:.2f}")
```

Output :

Confusion Matrix :

[[4]]

Accuracy :1.00

Precision :1.00

Recall :1.00

F1 Performance score :1.00

EX.NO :

DATE :

IMPLEMENT DECISION TREE CLASSIFIER

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import
confusion_matrix,accuracy_score,precision_score,recall_score,f1_score
```

Step 2:

```
[ ] data=pd.read_csv("DecisionTree_TargetIncome.csv")
print("the First 5 values of data is :\n")
data.head()
```

Output:

the First 5 values of data is :

	Age	Income	Gender	TargetVariable
0	25	50000	Male	0
1	45	64000	Female	1
2	35	58000	Female	0
3	50	72000	Male	1
4	23	48000	Male	0

Step 3:

```
[ ] data=pd.get_dummies(data, columns=['Gender'],drop_first=True)
```


Step 4:

```
[ ] print(data.columns)
```

Step 5:

```
[ ] X = data.drop('TargetVariable',axis=1)
    y = data['TargetVariable']
```

Step 6:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
    classifier = DecisionTreeClassifier(random_state=0)
    classifier.fit(X_train, y_train)
```

Output:

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

Step 7:

```
[ ] y_pred = classifier.predict(X_test)
```

Step 8:

```
[ ] cm = confusion_matrix(y_test, y_pred)
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
```

Step 9:

```
[ ] print(f"Confusion Matrix :\n", cm)
    print(f"Accuracy : {accuracy:.2f}")
    print(f"Precision : {precision:.2f}")
    print(f"Recall : {recall:.2f}")
    print(f"F1 score : {f1:.2f}")
```

Output :

Confusion Matrix :

[[0 0]

[1 3]]

Accuracy :0.75

Precision :1.00

Recall :0.75

F1 score :0.86

MLA-LAB
221421601012

EX.NO :

DATE :

IMPLEMENT RANDOM FOREST ALGORITHM

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
    from sklearn.model_selection import train_test_split
    from sklearn.ensemble import RandomForestClassifier
    from sklearn.metrics import accuracy_score, classification_report
    import warnings
    warnings.filterwarnings('ignore')
```

Step 2:

```
[ ] titanic_data=pd.read_csv('RandomForest_TitanicSurvival.csv')
    titanic_data.head()
```

Output:

	Passenger Id	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.25	-	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282	7.925	-	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803	53.1	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35	0	0	373450	8.05	-	S

Step 3:

```
[ ] titanic_data=titanic_data.dropna(subset=['Survived'])
```

Step 4:

```
[ ] X=titanic_data[['Pclass','Sex','Age','SibSp','Parch','Fare']]  
y=titanic_data['Survived']
```

Step 5:

```
[ ] X.loc[:, 'Sex']=X['Sex'].map({'female':0, 'male':1})  
X.loc[:, 'Age']=X['Age'].fillna(X['Age'].median())
```

Step 6:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
rf_classifier=RandomForestClassifier(n_estimators=100, random_state=42)  
rf_classifier.fit(X_train, y_train)
```

Output:

```
RandomForestClassifier  
RandomForestClassifier(random_state=42)
```

Step 7:

```
[ ] y_pred=rf_classifier.predict(X_test)
```

Step 8:

```
[ ] accuracy=accuracy_score(y_test,y_pred)  
classification_rep=classification_report(y_test,y_pred)
```

Step 9:

```
[ ] print(f"Accuracy : {accuracy:.2f}")  
print(f"Classification Report :\n", classification_rep)
```

Output :

Accuracy :0.80

Classification Report :

	<i>precision</i>	<i>recall</i>	<i>f1-score</i>	<i>support</i>
<i>0</i>	<i>0.82</i>	<i>0.85</i>	<i>0.83</i>	<i>105</i>
<i>1</i>	<i>0.77</i>	<i>0.73</i>	<i>0.75</i>	<i>74</i>
<i>accuracy</i>			<i>0.80</i>	<i>179</i>
<i>macro avg</i>	<i>0.79</i>	<i>0.79</i>	<i>0.79</i>	<i>179</i>
<i>weighted avg</i>	<i>0.80</i>	<i>0.80</i>	<i>0.80</i>	<i>179</i>

EX.NO :

DATE :

IMPLEMENT K-MEANS CLUSTERING ALGORITHM

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
import numpy as np
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
```

Step 2:

```
[ ] data=pd.read_csv("KMeans_FeatureClusters.csv")
print("the First 5 values of data is :\n")
data.head()
```

Output:

the First 5 values of data is :

	<i>Feature1</i>	<i>Feature2</i>
<i>0</i>	2.3	3.4
<i>1</i>	1.5	1.8
<i>2</i>	7.6	6.5
<i>3</i>	2.1	4.2
<i>4</i>	8.0	7.0

Step 3:

```
[ ] kmeans=KMeans(n_clusters=2,random_state=0)
    kmeans.fit(df)
```

Output:

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:870:

FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

warnings.warn(

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1382:

UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.

warnings.warn(

```
▼ KMeans
KMeans(n_clusters=2, random_state=0)
```

Step 4:

```
[ ] df['cluster']=kmeans.labels_
```

Step 5:

```
[ ] print(df.head())
```

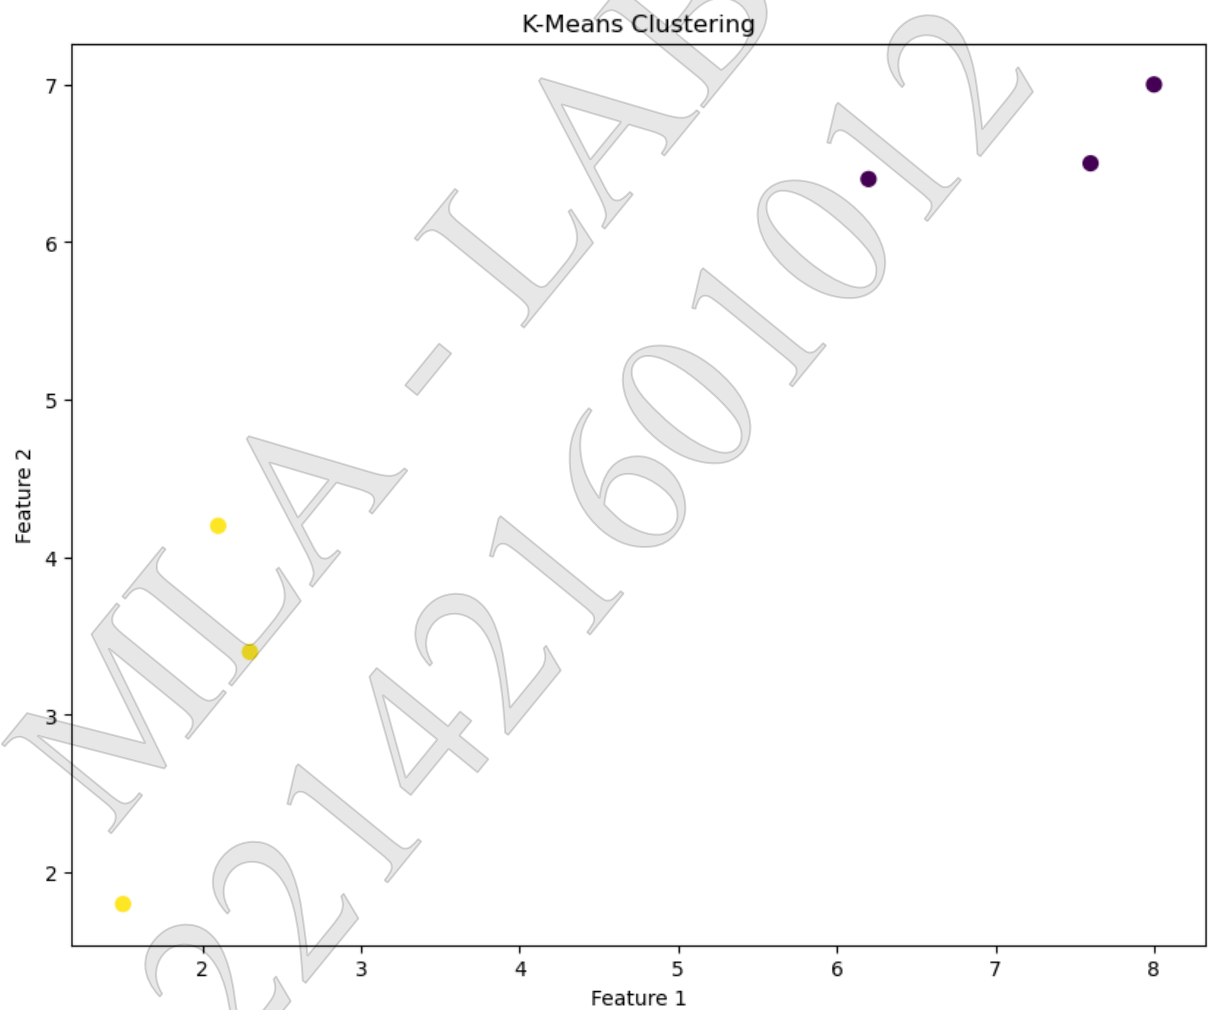
Output:

	<i>Feature1</i>	<i>Feature2</i>	<i>Cluster</i>
<i>0</i>	2.3	3.4	1
<i>1</i>	1.5	1.8	1
<i>2</i>	7.6	6.5	0
<i>3</i>	2.1	4.2	1
<i>4</i>	8.0	7.0	0

Step 6:

```
[ ] plt.figure(figsize=(10,8))  
    plt.scatter(df['Feature1'],df['Feature2'],c=df['cluster'],cmap='viridis',s=50)  
    plt.title('K-Means Clustering')  
    plt.xlabel('Feature 1')  
    plt.ylabel('Feature 2')  
    plt.show()
```

Output:



EX.NO :

DATE :

IMPLEMENT PRINCIPAL COMPONENT ANALYSIS (PCA)

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
```

Step 2:

```
[ ] data=pd.read_csv("PCA_CustomerSegmentation.csv")
print("the First 5 values of data is :\n")
data.head()
```

Output:

the First 5 values of data is :

	<i>Age</i>	<i>Annual Income</i>	<i>Spending Score</i>	<i>Number of Purchases</i>	<i>Cluster</i>
<i>0</i>	22	35000	45	15	3
<i>1</i>	28	45000	50	18	2
<i>2</i>	33	60000	55	22	1
<i>3</i>	38	70000	60	25	4
<i>4</i>	45	80000	65	30	2

Step 3:

```
[ ] Features=['Age','Annual Income','Spending Score','Number of Purchases']  
X=df[Features]
```

Step 4:

```
[ ] Scaler=StandardScaler()  
X_scaled=Scaler.fit_transform(X)
```

Step 5:

```
[ ] pca=PCA(n_components=2)  
X_pca=pca.fit_transform(X_scaled)
```

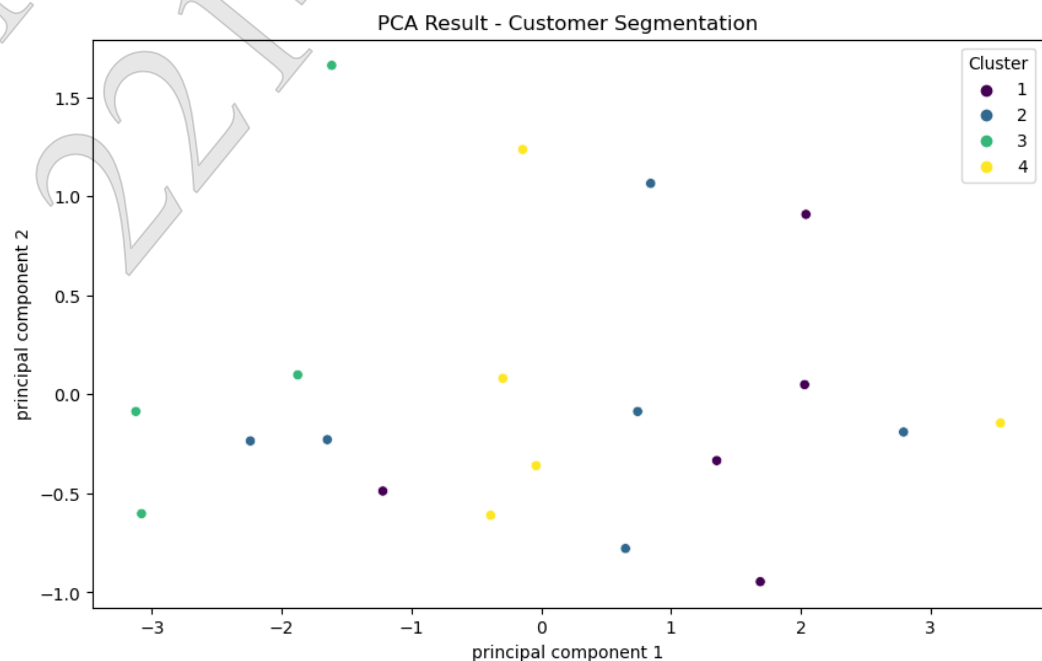
Step 6:

```
[ ] df_pca=pd.DataFrame(data=X_pca,columns=['pc1','pc2'])  
df_pca['Cluster']=df['Cluster']
```

Step 7:

```
[ ] plt.figure(figsize=(10,6))  
sns.scatterplot(x='pc1',y='pc2',hue='Cluster',data=df_pca,palette='viridis')  
plt.title('PCA Result - Customer Segmentation')  
plt.xlabel('principal component 1')  
plt.ylabel('principal component 2')  
plt.show()
```

Output :



Step 8:

```
[ ] explained_variance=pca.explained_variance_ratio_  
    print(f'Explained Variance by each Component: {explained_variance}')  
    print(f'Total Explained Variance : {np.sum(explained_variance)}')
```

Output :

Explained Variance by each Component: [0.87276806 0.1144451]

Total Explained Variance : 0.9872131621734076

MLA-LAB
221421601012

EX.NO :

DATE :

IMPLEMENT K-NEAREST NEIGHBORS (KNN) ALGORITHM

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
import plotly.graph_objects as go
import plotly.offline as pyoff
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

Step 2:

```
[ ] data=pd.read_csv("KNN_CellClassification.csv")
data.info()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 698 entries, 0 to 697
Data columns (total 11 columns):
#   Column  Non-Null Count  Dtype
---  -
0    1000025    698 non-null    int64
1     5         698 non-null    int64
2     1         698 non-null    int64
3    1.1        698 non-null    int64
4    1.2        698 non-null    int64
5     2         698 non-null    int64
6    1.3        698 non-null    object
7     3         698 non-null    int64
8    1.4        698 non-null    int64
9    1.5        698 non-null    int64
10   2.1        698 non-null    int64
dtypes: int64(10), object(1)
memory usage: 60.1+ KB
```

Step 3:

```
[ ] data.head()
```

Output:

	<i>1000025</i>	<i>5</i>	<i>1</i>	<i>1.1</i>	<i>1.2</i>	<i>2</i>	<i>1.3</i>	<i>3</i>	<i>1.4</i>	<i>1.5</i>	<i>2.1</i>
<i>0</i>	1002945	5	4	4	5	7	10	3	2	1	2
<i>1</i>	1015425	3	1	1	1	2	2	3	1	1	2
<i>2</i>	1016277	6	8	8	1	3	4	3	7	1	2
<i>3</i>	1017023	4	1	1	3	2	1	3	1	1	2
<i>4</i>	1017122	8	10	10	8	7	10	9	7	1	4


Step 4:

```
[ ] data.columns=['Id','Clump Thickness','uniformity of cell size','uniformity of cell  
shape','Marginal Adhesion','Single Epithelial cell size','Bare Nuclei','Bland  
Chromatin','Normal Nucleoli','Mitoses','Class']
```

Step 5:

```
[ ] target_balance=data['Class'].value_counts().reset_index()  
target_balance.columns=['Class','Count']  
target_class=go.Bar(  
    name="Target Balance",  
    x=target_balance['Class'].astype(str),  
    y=target_balance['Count']  
)  
fig=go.Figure(target_class)  
pyoff.iplot(fig)
```

Output:




Step 6:

```
[ ] beg_class_pat=data.loc[data['Class']==2]
mal_class_pat=data.loc[data['Class']==4]
Mith_10_beg=beg_class_pat['Mitoses'].value_counts().reset_index()
Mith_10_beg.columns=['Mitoses','Count']
Mith_10_mal=mal_class_pat['Mitoses'].value_counts().reset_index()
Mith_10_mal.columns=['Mitoses','Count']
```

Step 7:

```
[ ] fig=go.Figure(data=[
    go.Bar(name='Levels of Mitoses in Benign
    Group',x=Mith_10_beg['Mitoses'].astype(str),y=Mith_10_mal['Count']),
    go.Bar(name='Levels of Mitoses in Malignant
    Group',x=Mith_10_beg['Mitoses'].astype(str),y=Mith_10_mal['Count']))]
fig.update_layout(barmode='group')
fig.show()
```

Output:



Step 8:

```
[ ] x=data.drop(columns=['Id','Class'])
y=data['Class']
```

Step 9:

```
[ ] print("Unique Values in 'Bara Nuclei':",x['Bare Nuclei'].unique())
```

Output :

Unique Values in 'Bara Nuclei': ['10' '2' '4' '1' '3' '9' '7' '?' '5' '8' '6']

Step 10:

```
[ ] x['Bare Nuclei']=pd.to_numeric(x['Bare Nuclei'],errors='coerce')
```

Step 11:

```
[ ] x=x.fillna(x.median())
```

Step 12:

```
[ ] x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)
```

Step 13:

```
[ ] scaler=StandardScaler()
x_train_scaled=scaler.fit_transform(x_train)
x_test_scaled=scaler.transform(x_test)
```

Step 14:

```
[ ] knn=KNeighborsClassifier(n_neighbors=5)
knn.fit(x_train_scaled,y_train)
```

Output :

```
▼ KNeighborsClassifier
KNeighborsClassifier()
```

Step 15:

```
[ ] y_pred=knn.predict(x_test_scaled)
```

Step 16:

```
[ ] print("Accuracy: ",accuracy_score(y_test,y_pred))
    print()
    print("Confusion Matrix: \n",confusion_matrix(y_test,y_pred))
    print("Classification Report: \n",classification_report(y_test,y_pred))
```

Output :

Accuracy: 0.9714285714285714

Confusion Matrix:

[[131 4]

[2 73]]

Classification Report:

	precision	recall	f1-score	support
2	0.98	0.97	0.98	135
4	0.95	0.97	0.96	75
accuracy			0.97	210
macro avg	0.97	0.97	0.97	210
weighted avg	0.97	0.97	0.97	210

EX.NO :

DATE :

DEMONSTRATE WEKA TOOL

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
    from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import StandardScaler
    from sklearn.ensemble import RandomForestClassifier
    from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

Step 2:

```
[ ] data=pd.read_csv("Weka_CustomerChurn.csv")
    print("the First 5 values of data is :\n")
    data.head()
```

Output:

the First 5 values of data is :

	Customer ID	Gender	Senior Citizen	Partner	Dependents	Tenure	Phone Service	Monthly Charges	Total Charges	Churn
0	1	Male	0	Yes	No	1	No	29.85	29.85	No
1	2	Female	1	No	No	34	Yes	56.95	1889.50	No
2	3	Female	0	Yes	Yes	2	Yes	53.85	108.15	Yes
3	4	Male	0	No	Yes	45	Yes	42.30	1840.75	No
4	5	Female	1	No	No	2	No	70.70	151.65	Yes

Step 3:

```
[ ] X = data.drop(columns=['Churn'])  
    y = data['Churn'].apply(lambda x: 1 if x == 'Yes' else 0)
```

Step 4:

```
[ ] X = pd.get_dummies(X, drop_first=True)
```

Step 5:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 6:

```
[ ] scaler = StandardScaler()  
    X_train = scaler.fit_transform(X_train)  
    X_test = scaler.transform(X_test)
```

Step 7:

```
[ ] model = RandomForestClassifier(n_estimators=100, random_state=42)
```

Step 8:

```
[ ] model.fit(X_train, y_train)
```

Output:

```
▼ RandomForestClassifier  
RandomForestClassifier(random_state=42)
```

Step 9:

```
[ ] y_pred = model.predict(X_test)
```

Step 10:

```
[ ] accuracy = accuracy_score(y_test, y_pred)  
    conf_matrix = confusion_matrix(y_test, y_pred)  
    report = classification_report(y_test, y_pred)
```

Output :

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

Step 11:

```
[ ] print(f'Accuracy: {accuracy}')  
    print('Confusion Matrix:')  
    print(conf_matrix)  
    print('Classification Report:')  
    print(report)
```

Output :

Accuracy: 0.5

Confusion Matrix:

[[0 1]

[0 1]]

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.50	1.00	0.67	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

EX.NO :

DATE :

BUILD AN UNSUPERVISED MODEL USING [Scikit-Learn and PyTorch]

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
    from sklearn.datasets import load_iris
    from sklearn.preprocessing import StandardScaler
```

Step 2:

```
[ ] iris=load_iris()
    X=pd.DataFrame(iris.data,columns=iris.feature_names)
```

Step 3:

```
[ ] Scaler=StandardScaler()
    X_Scaled=Scaler.fit_transform(X)
```

Step 4:

```
[ ] from sklearn.cluster import KMeans
    from sklearn.metrics import silhouette_score
    kmeans=KMeans(n_clusters=3,random_state=42)
    kmeans_labels=kmeans.fit_predict(X_Scaled)
```

Output :

```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the
value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1382:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there
are less chunks than available threads. You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
  warnings.warn(
```

Step 5:

```
[ ] kmeans_silhouette=silhouette_score(X_Scaled, kmeans_labels)
    print(f'K-Means Silhouette Score: {kmeans_silhouette}')
```

Output :

K-Means Silhouette Score: 0.45994823920518635

Step 6:

```
[ ] import torch
    import torch.nn as nn
    import torch.optim as optim
    from torch.utils.data import DataLoader, TensorDataset
    X_tensor=torch.tensor(X_Scaled, dtype=torch.float32)
```

Step 7:

```
[ ] data_loader=DataLoader(TensorDataset(X_tensor),batch_size=16,shuffle=True)
```

Step 8:

```
[ ] class Autoencoder(nn.Module):
    def __init__(self):
        super(Autoencoder, self).__init__()
        self.encoder=nn.Sequential(
            nn.Linear(4,2),
            nn.ReLU()
        )
        self.decoder=nn.Sequential(
            nn.Linear(2,4),
            nn.ReLU()
        )
    def forward(self,x):
        encoded=self.encoder(x)
        decoded=self.decoder(encoded)
        return decoded
```

Step 9:

```
[ ] autoencoder=Autoencoder()  
    criterion=nn.MSELoss()  
    optimizer=optim.Adam(autoencoder.parameters(),lr=0.01)
```

Step 10:

```
[ ] num_epochs=100  
    for epoch in range(num_epochs):  
        for data in data_loader:  
            inputs, =data  
            optimizer.zero_grad()  
            outputs=autoencoder(inputs)  
            loss=criterion(outputs, inputs)  
            loss.backward()  
            optimizer.step()
```

Step 11:

```
[ ] with torch.no_grad():  
    encoded_data=autoencoder.encoder(X_tensor).numpy()
```

Step 12:

```
[ ] encoded_kmeans=KMeans(n_clusters=3,random_state=42)  
    encoded_kmeans_labels=encoded_kmeans.fit_predict(encoded_data)
```

Output :

```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:870:  
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set  
the value of `n_init` explicitly to suppress the warning  
  warnings.warn(  
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1382:  
UserWarning: KMeans is known to have a memory leak on Windows with MKL,  
when there are less chunks than available threads. You can avoid it by setting the  
environment variable OMP_NUM_THREADS=1.  
  warnings.warn(  

```

Step 13:

```
[ ] encoded_silhouette = silhouette_score(X_Scaled, encoded_kmeans_labels)
    print(f'Autoencoder + kmeans silhouette score : {encoded_silhouette}')
```

Output :

Autoencoder + kmeans silhouette score : 0.302958422921628

MLA-LAB
221421601012

EX.NO :

DATE :

IMPLEMENT LOGISTIC REGRESSION

SOURCE CODE :

Step 1:

```
[ ] import pandas as pd
    from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import StandardScaler, LabelEncoder
    from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Step 2:

```
[ ] data=pd.read_csv("LogisticRegression_HeartDisease.csv")
    print("the First 5 values of data is : \n")
    data.head()
```

Output:

the First 5 values of data is :

	Age	CholesterolLevel	BloodPressure	SmokingStatus	HeartDisease
0	68	213.510222	119.336301	No	No
1	58	209.232228	109.821694	Yes	No
2	44	194.023226	154.003458	No	No
3	72	152.111973	162.855094	Yes	No
4	37	179.826361	121.379939	No	Yes

Step 3:

```
[ ] label_encoder = LabelEncoder()
    data['SmokingStatus'] = label_encoder.fit_transform(data['SmokingStatus'])
    data['HeartDisease'] = label_encoder.fit_transform(data['HeartDisease'])
```


Step 4:

```
[ ] X = data[['Age', 'CholesterolLevel', 'BloodPressure', 'SmokingStatus']]
    X.head()
```

Output:

	<i>Age</i>	<i>CholesterolLevel</i>	<i>BloodPressure</i>	<i>SmokingStatus</i>
0	68	213.510222	119.336301	0
1	58	209.232228	109.821694	1
2	44	194.023226	154.003458	0
3	72	152.111973	162.855094	1
4	37	179.826361	121.379939	0

Step 5:

```
[ ] y = data[['HeartDisease']]
    y.head()
```

Output:

	<i>HeartDisease</i>
0	0
1	0
2	0
3	0
4	1

Step 6:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
```

Step 7:

```
[ ] scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)
```

Step 8:

```
[ ] model = LogisticRegression()  
    model.fit(X_train, y_train)
```

Output:

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

```
▼ LogisticRegression  
LogisticRegression()
```

Step 9:

```
[ ] y_pred = model.predict(X_test)
```

Step 10:

```
[ ] accuracy = accuracy_score(y_test, y_pred)  
    conf_matrix = confusion_matrix(y_test, y_pred)  
    class_report = classification_report(y_test, y_pred)
```

Step 11:

```
[ ] print(f'Accuracy: {accuracy:.2f}')  
    print('Confusion Matrix:')  
    print(conf_matrix)  
    print('Classification Report:')  
    print(class_report)
```

Output:

Accuracy: 0.60

Confusion Matrix:

[[7 3]

[5 5]]

Classification Report:

	precision	recall	f1-score	support
0	0.58	0.70	0.64	10
1	0.62	0.50	0.56	10
accuracy			0.60	20
macro avg	0.60	0.60	0.60	20
weighted avg	0.60	0.60	0.60	20

EX.NO :

DATE :

IMPLEMENT NEURAL NETWORK MODEL

SOURCE CODE :

Step 1:

```
[ ] import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

Step 2:

```
[ ] df=pd.read_csv("NeuralNetwork_PredictiveMaintenance.csv")
print("the First 5 values of data is :\n")
df.head()
```

Output:

The First 5 values of data is :

	sensor_1	sensor_2	sensor_3	sensor_4	sensor_5	sensor_6	sensor_7	sensor_8	sensor_9	sensor_10	failure
0	0.374540	0.950714	0.731994	0.598658	0.156019	0.155995	0.058084	0.866176	0.601115	0.708073	0
1	0.020584	0.969910	0.832443	0.212339	0.181825	0.183405	0.304242	0.524756	0.431945	0.291229	1
2	0.611853	0.139494	0.292145	0.366362	0.456070	0.785176	0.199674	0.514234	0.592415	0.046450	1
3	0.607545	0.170524	0.065052	0.948886	0.965632	0.808397	0.304614	0.097672	0.684233	0.440152	1
4	0.122038	0.495177	0.034389	0.909320	0.258780	0.662522	0.311711	0.520068	0.546710	0.184854	0

Step 3:

```
[ ] X = df.drop('failure', axis=1).values # Features (sensor data)
y = df['failure'].values # Target (failure labels)
```

Step 4:

```
[ ] scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)
```

Step 5:

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.15,
    random_state=42)

    X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.15,
    random_state=42)
```

Step 6:

```
[ ] model = Sequential()
    model.add(Dense(64, input_dim=X.shape[1], activation='relu'))
    model.add(Dense(32, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))
```

Output:

C:\ProgramData\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Step 7:

```
[ ] model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

Step 8:

```
[ ] # Train the model with fewer epochs
    history = model.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=10,
    batch_size=32)

    # Evaluate the model on the test set
    test_loss, test_accuracy = model.evaluate(X_test, y_test)
    print(f"Test Loss: {test_loss}")
    print(f"Test Accuracy: {test_accuracy}")

    # Predict and generate classification report
    y_pred = (model.predict(X_test) > 0.5).astype("int32")
    print(classification_report(y_test, y_pred))
```

Output:


```
Epoch 1/10
23/23 ————— 0s 4ms/step - accuracy: 0.7976 - loss: 0.5117 -
val_accuracy: 0.5156 - val_loss: 0.7832
Epoch 2/10
23/23 ————— 0s 2ms/step - accuracy: 0.7784 - loss: 0.5235 -
val_accuracy: 0.5156 - val_loss: 0.7902
Epoch 3/10
23/23 ————— 0s 3ms/step - accuracy: 0.7891 - loss: 0.5068 -
val_accuracy: 0.5000 - val_loss: 0.7948
Epoch 4/10
23/23 ————— 0s 3ms/step - accuracy: 0.7687 - loss: 0.5194 -
val_accuracy: 0.4922 - val_loss: 0.8020
Epoch 5/10
23/23 ————— 0s 2ms/step - accuracy: 0.7619 - loss: 0.5247 -
val_accuracy: 0.5312 - val_loss: 0.8121
Epoch 6/10
23/23 ————— 0s 2ms/step - accuracy: 0.7967 - loss: 0.5169 -
val_accuracy: 0.5234 - val_loss: 0.8118
Epoch 7/10
23/23 ————— 0s 2ms/step - accuracy: 0.7831 - loss: 0.5106 -
val_accuracy: 0.5000 - val_loss: 0.8141
Epoch 8/10
23/23 ————— 0s 3ms/step - accuracy: 0.7767 - loss: 0.5078 -
val_accuracy: 0.5391 - val_loss: 0.8291
Epoch 9/10
23/23 ————— 0s 2ms/step - accuracy: 0.7946 - loss: 0.4805 -
val_accuracy: 0.4922 - val_loss: 0.8261
Epoch 10/10
23/23 ————— 0s 2ms/step - accuracy: 0.8199 - loss: 0.4780 -
val_accuracy: 0.5156 - val_loss: 0.8287
5/5 ————— 0s 2ms/step - accuracy: 0.4935 - loss: 0.7497
Test Loss: 0.7296823263168335
Test Accuracy: 0.5
5/5 ————— 0s 2ms/step
```

	precision	recall	f1-score	support
0	0.48	0.44	0.46	73
1	0.51	0.56	0.53	77
accuracy			0.50	150
macro avg	0.50	0.50	0.50	150
weighted avg	0.50	0.50	0.50	150

Step 9:

```
[ ] # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Failure',
'Failure'], yticklabels=['No Failure', 'Failure'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```


Output:



Step 10:

```
[ ] # Plot learning curves for loss
plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Learning Curve: Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Output:



Step 11:

```
[ ] # Plot learning curves for accuracy
plt.figure(figsize=(10, 5))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Learning Curve: Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

Output:

